

Auto-Encoder Reconstruction Cost for Anomalies Detection scenarios

**Comparing Auto-Encoder approach for Anomalies Detection to
traditional Upsampling + Classification techniques**

Dataset

Credit Card Fraud Detection


Anonymized credit card transactions labeled as fraudulent or genuine

<https://www.kaggle.com/mlg-ulb/creditcardfraud>

2 days real European credit cards transactions: **492 frauds out of 284,807 (0.172%)**


RAW

- 28 Numerical input variables which are the result of a PCA transformation to anonymize customer info
- Transaction Amount
- Time in seconds from T0
- Label Class: 0/1 (1 is the anomaly)



Transform

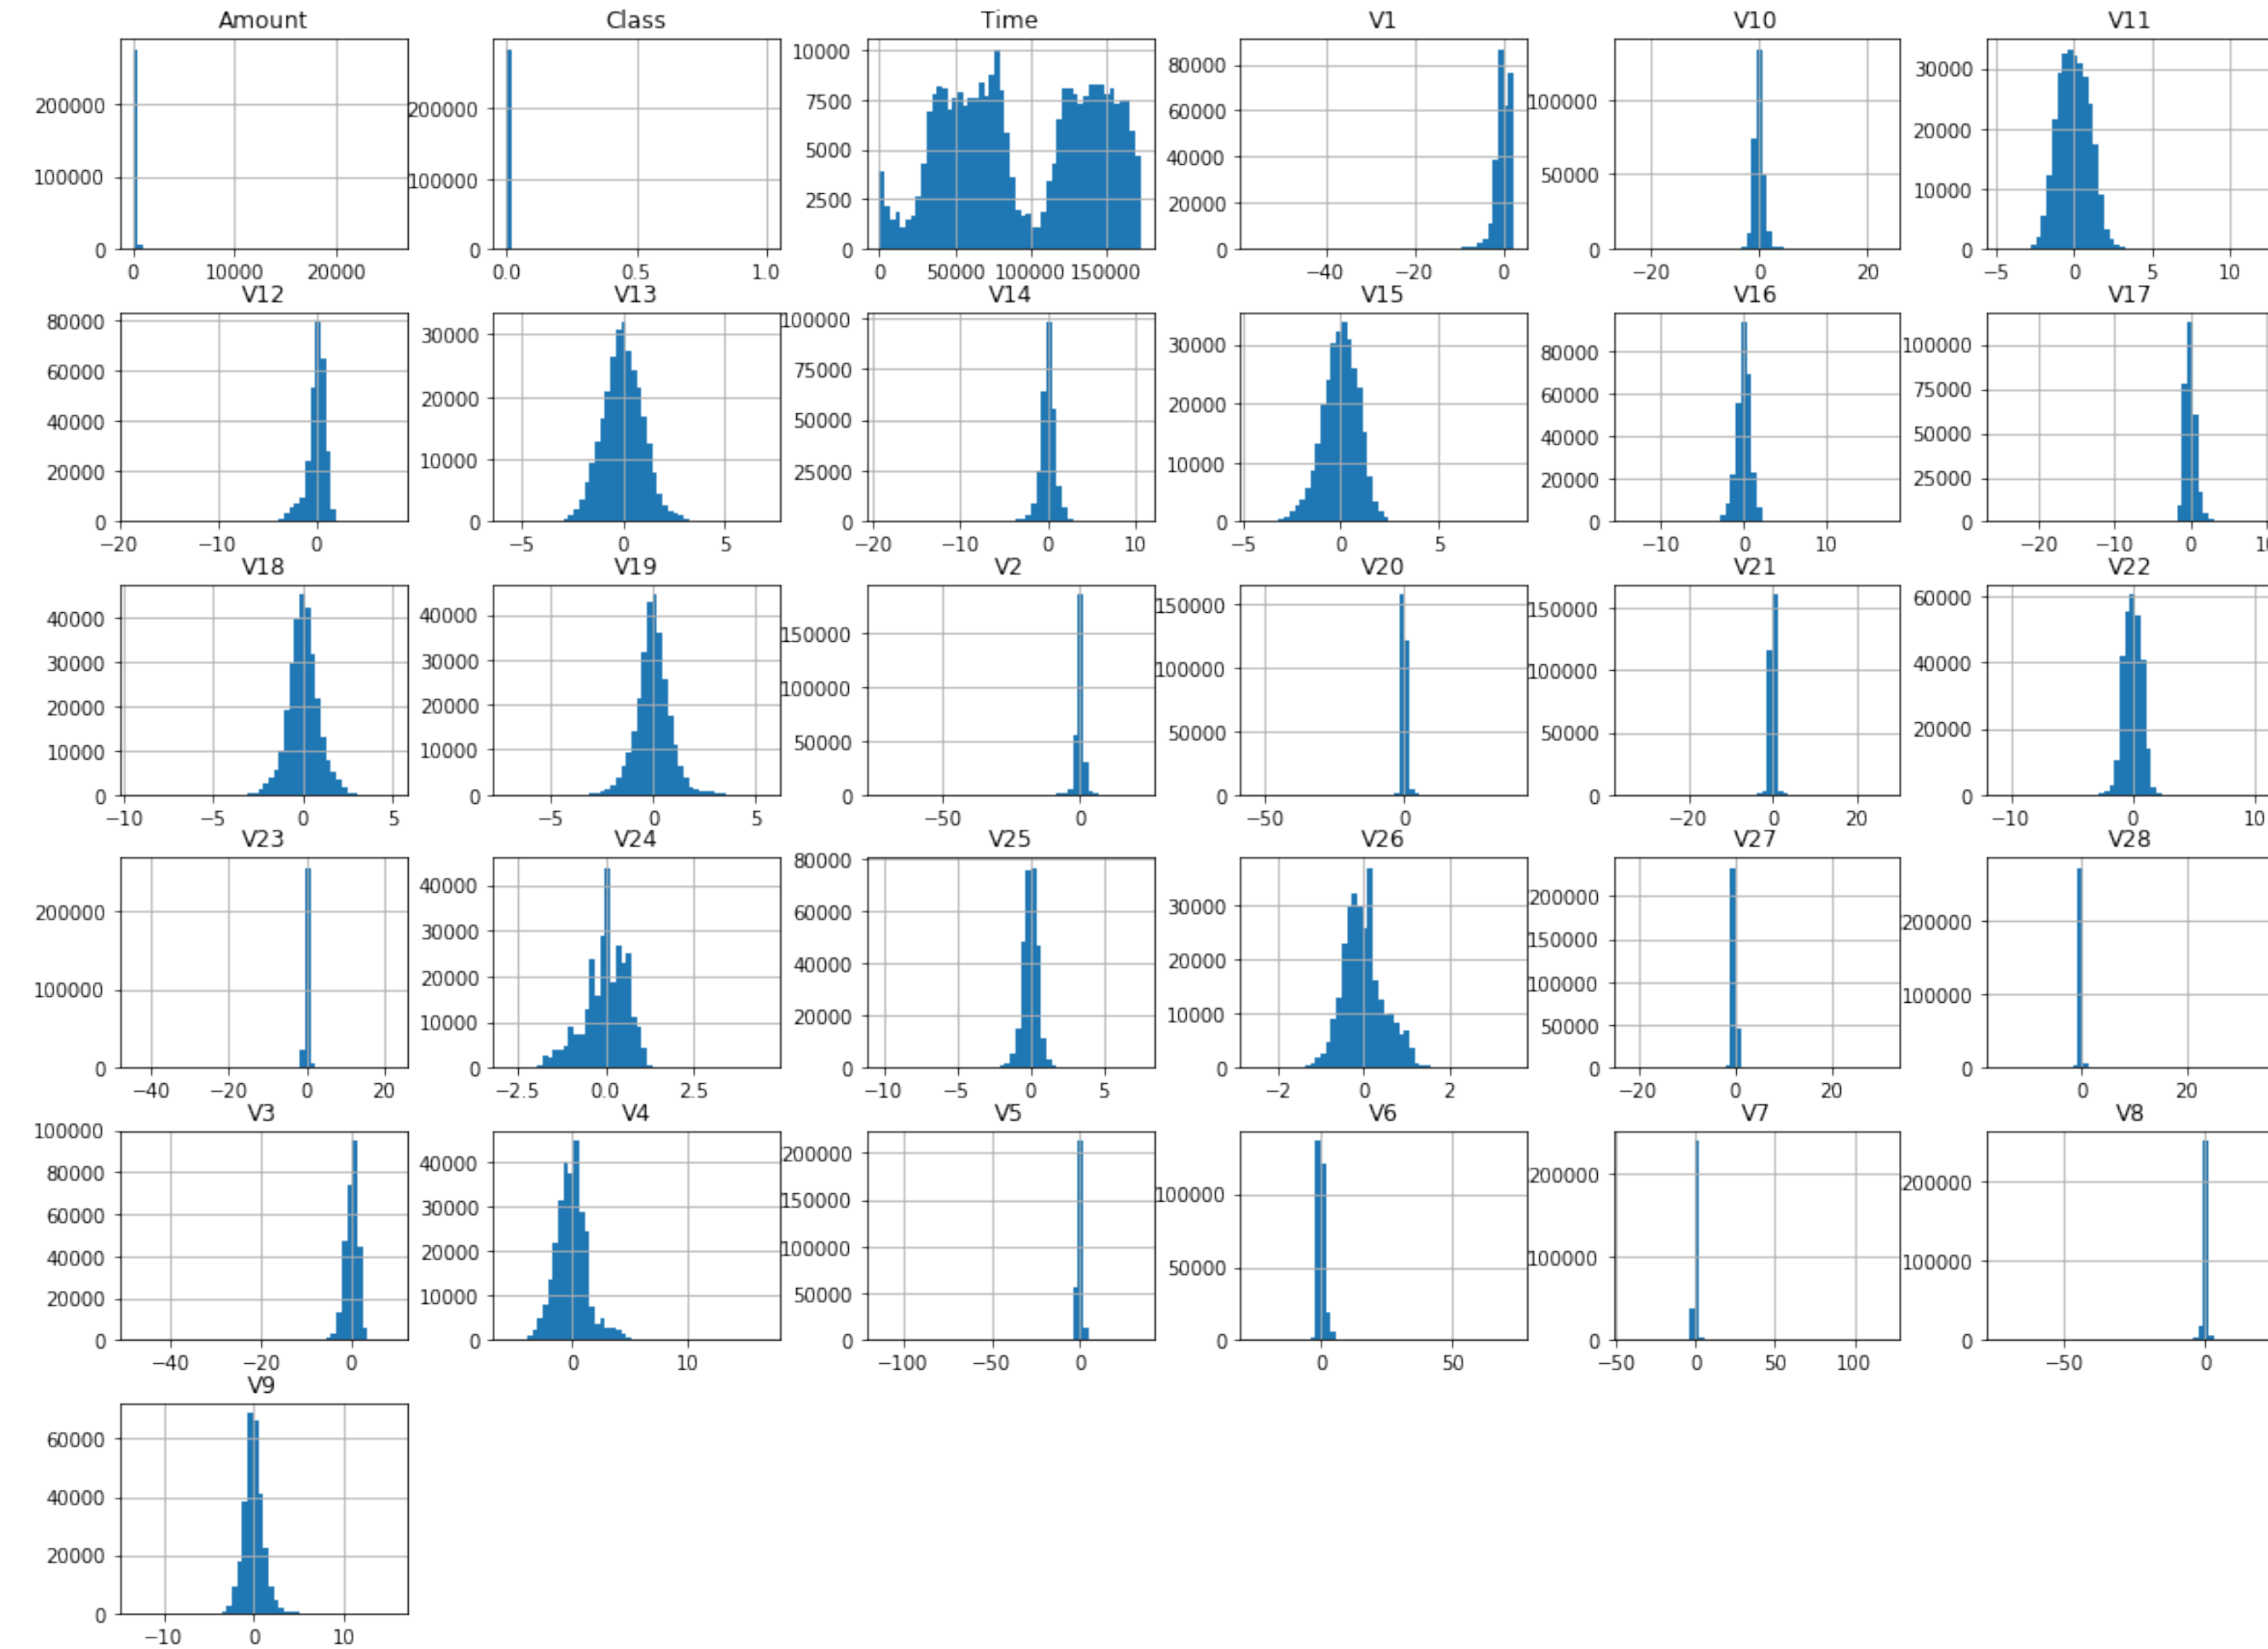
- 29 Numerical Features
 - Original 28 from PCA
 - Transaction Amount (StandardScaler)
- 1 Categorical Feature
 - Time of the transaction using 24 Hour bins




Multi-Input

- 29 Numerical Features -> Tensor Length 29
- 1 Categorical Feature values 0-24 —> Embedding Tensor 12


Plot



Time in 48 total hours



Time in 24 hours bin



Training Approach

Classifier

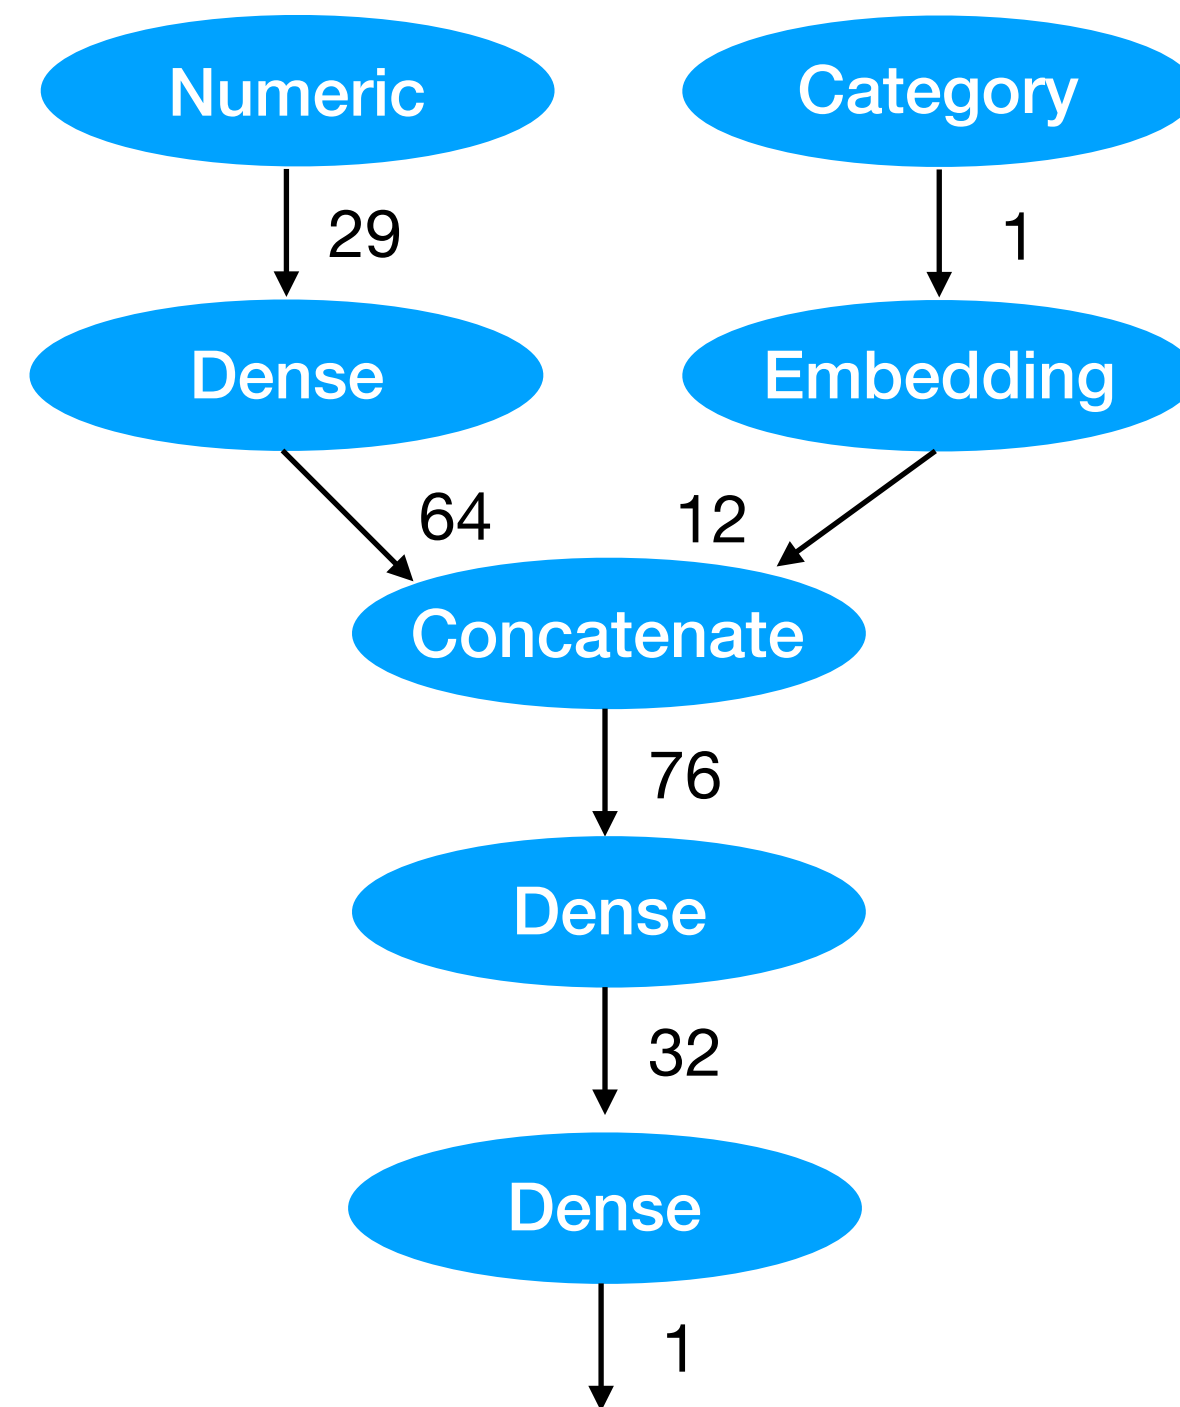
- Feed 80% of both anomaly (+) and non-anomaly (-) transactions
- Upscale the anomaly transactions in the train dataset
- Handle Multi-Input for Numeric features and Categories
- Use Embedding for Categories
- Build a classic DNN classifier

Auto-Encoder

- Feed 80% of only the non-anomaly (-) transactions
- Handle Multi-Input for Numeric features and Categories
- Use Embedding for Categories
- Use an Auto-Encoder to learn compression/decompression on non-anomaly transactions
- Use a custom layer to calculate the reconstruction cost
- Train to reduce this cost to 0




Classifier



Models





Auto-Encoder



Reconstruction Cost
$$\sum (\text{Encoder_Input} - \text{Decoder_Output})^2$$

Classifier

accuracy: 0.9839015484006882
precision: [0.99980346 0.08853119]
recall: [0.98406697 **0.88888889**]
fscore: [0.9918728 0.1610247]
support: [56863 99]





Results

Feed 20% **+** **-**
SCORING

Auto-Encoder (THRESHOLD = 0.17)


accuracy: 0.9200344089041818
precision: [0.99969431 0.01795759]
recall: [0.92017656 **0.83838384**]
fscore: [0.95828869 0.03516204]
support: [56863 99]



Interactivity test

THRESHOLD  0.17

Feed 80%



SCORING

TRAINED NEGATIVE

Greater Than Threshold: 13932 over 227452

Accuracy: 0.94


Mean: 0.05370073765516281

Std: 0.09088834375143051

Min: 5.205785419093445e-05

Max: 3.160778045654297

Feed 20%



SCORING

NEW NEGATIVE

Greater Than Threshold: 4809 over 56863

Accuracy: 0.92


Mean: 0.07014649361371994

Std: 0.1701817512512207

Min: 8.863704715622589e-05

Max: 26.871849060058594

Feed 100%



SCORING

NEW POSITIVE

Greater Than Threshold: 416 over 492

Accuracy: 0.85


Mean: 3.5444180965423584

Std: 7.813464164733887

Min: 0.001192888943478465

Max: 55.32634735107422

Features Interpretability



Source Code

<https://github.com/JacopoMangiavacchi/AutoEncoderForAnomalyDetection>